

Mastering DeepScale R

Build and deploy AI Models with ollama



• what is DeepScaleR?

- Fine-tuned version of DeepSeek-R1-Distilled-Qwen-1.5B
- Specifically designed to surpass open AI's o1 preview model in mathematical reasoning tasks.
- High accuracy and efficiency
- Extremely lightweight yet powerful model for AI-driven computations.

• why choose DeepScaleR?

- lightweight (1.5B parameters) yet powerful
- Outperforms larger models in structured math evaluations.
- Runs locally with ollama (No API costs!)
- Supports fine-tuning & Customization
- low-latency inference (ideal for real-time applications)

• Overview of DeepScale R's Capabilities

1. Distilled from DeepSeek-R1 for Efficiency.

- DeepSeek-R1 is a family of open-source LLMs with strong performance in text generation and reasoning.
- DeepScaleR refines this architecture to be optimized for math-heavy and structured problem-solving tasks.
- It retains the speed and efficiency of DeepSeek-R1 while enhancing accuracy for numerical and symbolic reasoning.

2. Advanced Mathematical Capabilities

- Optimized for symbolic, algebraic, and numerical tasks, including :
 - Arithmetic and Algebra { $+$, $-$, $*$, $/$ and algebraic equations }
 - Linear Algebra { matrix operations, determinants and eigen values }
 - Calculus { Derivatives, Integrals, limits and series expansion }
 - Probability and statistics { Bayesian analysis, hypothesis testing }
 - Logical and Symbolic Reasoning { Theorem proving, formal logic }

3. Small yet powerful (1.5B Parameters)

- DeepScaleR is significantly smaller than most mainstream LLMs yet delivers comparable (or better) performance in its target areas.
- Designed for low-latency inference, making it ideal for edge devices, personal AI assistants, and embedded AI Systems.
- Supports quantization techniques to further reduce memory footprint while retaining accuracy.

4. Optimized for Deployment

- Compatible with ollama, enabling fast and easy local inference.
- Can run on CPUs and GPUs, making it more accessible than larger models.
- Supports fine-tuning with LoRA and QLoRA, allowing easy customization for specific domains

Comparison :

DeepScaleR vs. OpenAI's o1 - Preview

Feature	DeepScaleR (1.5B)	Open AI's o1 - preview
Model Size	1.5B Parameters	9B + Parameters
Performance on Math tasks	Better in structured math evaluations (algebra, calculus, linear algebra)	Strong in general problem-solving but weaker in structured math.
Inference speed	Fast (Optimized for low-latency inference)	Slower due to larger model size
Resource Requirement	lightweight, runs on CPU and GPU	Heavier, requires GPU acceleration
Fine-Tuning capability	LoRA/QLoRA supported for domain specific improvements	Limited fine-tuning options
Deployment Ease	Runs with Ollama, easy local setup	Requires cloud-based APIs
Licence & Accessibility	open-source & locally deployable	Closed-source, API-only

Key takeaway

DeepScaleR outperforms OpenAI-o1-preview in mathematical reasoning, has a much smaller footprint and is ideal for local AI deployments.

• Performance On Math Evaluations

• Benchmarking against other models

– DeepScaleR was tested on multiple math reasoning Datasets:

- GSMBK (Grade School Math 8K) → arithmetic & word problems
- MATH Dataset → High school and Olympiad level math.
- Hendrycks MATH → complex math reasoning
- APPS (AI problem solving) → Algorithm and problem solving tasks.

Model	GSMBK (Accuracy %)	MATH Dataset	Inference speed
DeepScaler (1.5 B)	85 %	78 %	Fast (CPU + GPU)
Open AI's o1-preview	80 %	70 %	Slower (GPU)
GPT-3.5-turbo	82 %	74 %	Medium
Mistral - 7 B	79 %	71 %	slow (GPU)

– Key Insights from Performance Evaluations

- DeepScaleR outperforms OpenAI's o1-preview on structured math datasets.
- Achieves near GPT 3.5 levels of accuracy in problem-solving with only 1.5B parameters.
- Much faster inference speed compared to larger models

• what is Ollama ?

- open-source framework that enables easy local deployment and inference of large language Models.
- Designed for efficiency, simplicity, and fast API model execution on personal computers, workstations, and servers.
- Download, run, and interact with models like DeepScaleR, Mistral, Llama, and more without needing complex cloud based APIs and GPU clusters.

• Introduction to the Ollama Framework

• What Does Ollama do ?

- Runs LLMs locally without requiring cloud-based APIs.
- Supports optimized models like DeepScaleR, Mistral, Qwen, Llama and other open-source LLMs.
- Provides a simple interface to interact with models via CLI, API, or Python SDK
- Works efficiently on both CPU and GPU, making AI accessible on a variety of machines.

• Key Features of ollama

- Local AI Execution - No external API calls & runs entirely on our machine.
- Lightweight & Optimized - uses quantized models for efficiency.
- Model Management - can easily pull, update and switch between models.
- Build-in API - offers a simple way to integrate AI into applications.
- Secure & private - No data leaves your local environment.

• How Does ollama Work ?

- Ollama provides a unified way to interact with different LLMs using a simple Command Line Interface [CLI] or API
- Download & Run Models
- Interact via API

• why use Ollama for AI Inference?

- **No More Expensive API Calls:** Instead of relaying on cloud based APIs, like openAI or anthropic that charge per request. Ollama allows to run AI models locally, saving cost and ensuring unlimited usage.
- **Faster Response Times:** It eliminates the latency issue of cloud hosted models, making it ideal for real time applications like chatbots, assistants, and AI driven automation.
- **Works Offline:** It does not require an internet connection, making it perfect for:
 - Edge AI Applications (AI running on local devices)
 - Privacy - Focused AI (no cloud data transmission)
 - AI for secure Environments (finance, defence, healthcare, etc)
- **Supports Multiple Models:** Ollama supports multiple LLM architectures, so can switch models on demand.
 - `ollama run deepScaleR` (for math)
 - `ollama run mistral` (for general reasoning)
 - `ollama run qwen` (for multilingual tasks)
- **Optimized for CPUs and GPUs**
 - Runs efficiently on consumer grade hardware
 - Supports quantization (like GbML, GPTQ) for lower memory usage and better speed.
 - No need for high-end GPUs, but scales well with them

• How Ollama Simplifies LLM Deployment

• Easy Installation & Setup

- Unlike traditional AI Development, which requires python environments, Cuda, and dependency management.

• Install ollama (Mac / Linux / Windows WSL)

- Go to [Ollama.com](https://ollama.com)
- Download on your machine

• Pull a model

- Go to terminal
- Check version: `ollama -v`
- Pull model: `ollama pull deepscaler`

• Run a model locally

- Run model: `ollama run deepscaler`
- Chat with it
- Exit model: `/bye`

• Built-in Model Serving

- Ollama includes a local API server, so you don't need to set up Flask, FastAPI, or other servers separately.

- Seamless Integration with Applications:

- Python SDK
- JavaScript API
- Dockerized AI Apps

• Why use Ollama for DeepScaleR?

- No Cloud Costs: Runs entirely on your local machine
- Fast AI Inference: low latency, ideal for real time AI
- Easy Deployment: CLI and API make model serving effortless
- Supports multiple models: Switching between AI models easily
- Private & Secure: No data leaves your local system

Hands - On Project

1. Setting Up the Environment

- a. Install ollama
- b. Download and Run DeepScaleR locally
- c. Running inference with DeepScaleR

2. Project 1: AI - Powered Math Solver

- Use DeepScaleR to build an AI-powered math solver that can handle complex math problems such as algebra, calculus, and linear algebra.
- Step by step:
 - Using DeepScaleR for solving mathematical equations
 - Implementing a web-based math assistant using Python & Gradio.

3. Project 2: AI Chatbot using DeepScaleR

- Handle general-purpose conversations, and we will deploy it as an API while ensuring low-latency responses
- Step 1: Building a chatbot for General - Purpose Conversation
- Step 2: Deploying DeepScaleR in an API
- Step 3: Building a web-based Chatbot UI with Gradio